Divide-and-Conquer



Divide-and-conquer.

- Break up problem into several parts.
- Solve each part recursively.
- Combine solutions to sub-problems into overall solution.

Most common usage.

- Break up problem of size n into two equal parts of size ½n.
- Solve two parts recursively.
- Combine two solutions into overall solution in linear time.

Consequence.

- Brute force: n².
- Divide-and-conquer: n log n.

Divide et impera.

Veni, vidi, vici.

- Julius Caesar





5.3 Counting Inversions



Counting Inversions



Music site tries to match your song preferences with others.

- You rank n songs.
- Music site consults database to find people with similar tastes.

Similarity metric: number of inversions between two rankings.

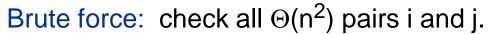
- . My rank: 1, 2, ..., n.
- Your rank: a₁, a₂, ..., a_n.
- Songs i and j inverted if i < j, but $a_i > a_j$.

Songs

	Α	В	С	D	Е				
Me	1	2	3	4	5				
You	1	3	4	2	5				

Inversions

3-2, 4-2



Applications



Applications.

- Voting theory.
- Collaborative filtering.
- Measuring the "sortedness" of an array.
- Sensitivity analysis of Google's ranking function.
- Rank aggregation for meta-searching on the Web.
- Nonparametric statistics (e.g., Kendall's Tau distance).
- $_{n}$ K(L,M)=

$$|\{(i,j)|i < j \land [((L(i) < L(j)) \land (M(i) > M(j))) \\ \lor ((L(i) > L(j)) \land (M(i) < M(j))) \]\}|$$





Divide-and-conquer.

1	5	4	8	10	2	6	9	12	11	3	7

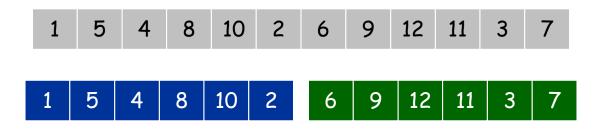




Divide: O(1).

Divide-and-conquer.

Divide: separate list into two pieces.







Divide-and-conquer.

- Divide: separate list into two pieces.
- Conquer: recursively count inversions in each half.



5 blue-blue inversions

8 green-green inversions

5-4, 5-2, 4-2, 8-2, 10-2

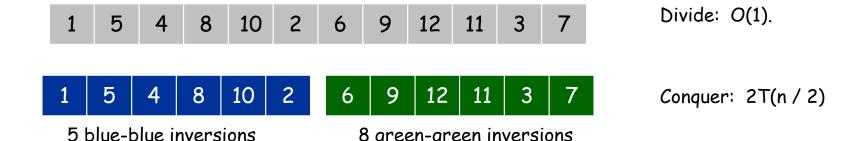
6-3, 9-3, 9-7, 12-3, 12-7, 12-11, 11-3, 11-7





Divide-and-conquer.

- Divide: separate list into two pieces.
- Conquer: recursively count inversions in each half.
- Combine: count inversions where a_i and a_i are in different halves, and return sum of three quantities.



8 green-green inversions

9 blue-green inversions 5-3, 4-3, 8-6, 8-3, 8-7, 10-6, 10-9, 10-3, 10-7

Combine: ???



Total = 5 + 8 + 9 = 22.

Counting Inversions: Combine



Combine: count blue-green inversions

- Assume each half is sorted.
- Count inversions where a_i and a_i are in different halves.
- Merge two sorted halves into sorted whole.

to maintain sorted invariant

25



13 blue-green inversions: 6 + 3 + 2 + 2 + 0 + 0

Count: O(n)

2 3 7 10

.0 11

14

16 1

18

19 23

Merge: O(n)

 $T(n) \leq T\left(\left\lfloor n/2\right\rfloor\right) + T\left(\left\lceil n/2\right\rceil\right) + O(n) \implies T(n) = O(n\log n)$

Counting Inversions: Implementation



Pre-condition. [Merge-and-Count] A and B are sorted. Post-condition. [Sort-and-Count] L is sorted.

```
Sort-and-Count(L) {
   if list L has one element
      return 0 and the list L

   Divide the list into two halves A and B
   (r<sub>A</sub>, A) ← Sort-and-Count(A)
   (r<sub>B</sub>, B) ← Sort-and-Count(B)
   (r , L) ← Merge-and-Count(A, B)

return r = r<sub>A</sub> + r<sub>B</sub> + r and the sorted list L
}
```





Merge-and-Count(A,B)

Maintain a Current pointer into each list, initialized to point to the front elements

Maintain a variable *Count* for the number of inversions, initialized to 0

While both lists are nonempty:

Let a_i and b_j be the elements pointed to by the *Current* pointer Append the smaller of these two to the output list If b_i is the smaller element then

Increment Count by the number of elements remaining in A Endif

Advance the Current pointer in the list from which the smaller element was selected.

EndWhile





```
Sort-and-Count(L)
```

If the list has one element then there are no inversions

Else

Divide the list into two halves:

A contains the first $\lceil n/2 \rceil$ elements

B contains the remaining $\lfloor n/2 \rfloor$ elements

 $(r_A, A) = Sort-and-Count(A)$

 $(r_B, B) = Sort-and-Count(B)$

(r, L) = Merge-and-Count(A, B)

Endif

Return $r = r_A + r_B + r$, and the sorted list L





The second secon

Closest pair. Given n points in the plane, find a pair with smallest Euclidean distance between them.

Fundamental geometric primitive.

- Graphics, computer vision, geographic information systems, molecular modeling, air traffic control.
- Special case of nearest neighbor, Euclidean MST, Voronoi.

fast closest pair inspired fast algorithms for these problems

Brute force. Check all pairs of points p and q with $\Theta(n^2)$ comparisons.

1-D version. O(n log n) easy if points are on a line.

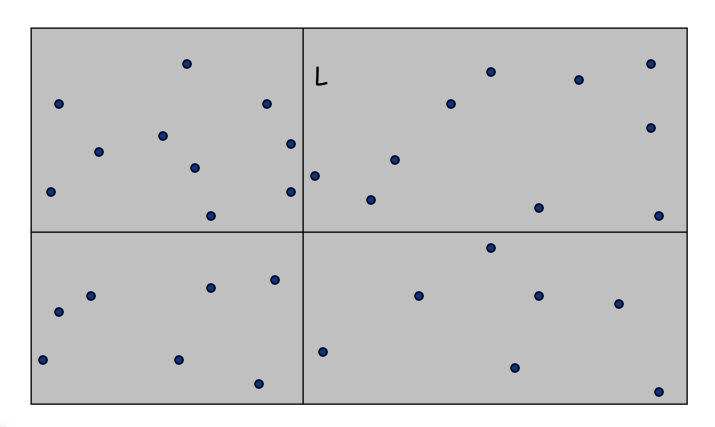
Assumption. No two points have same x coordinate.

to make presentation cleaner

Closest Pair of Points: First Attempt

A STATE OF THE STA

Divide. Sub-divide region into 4 quadrants.



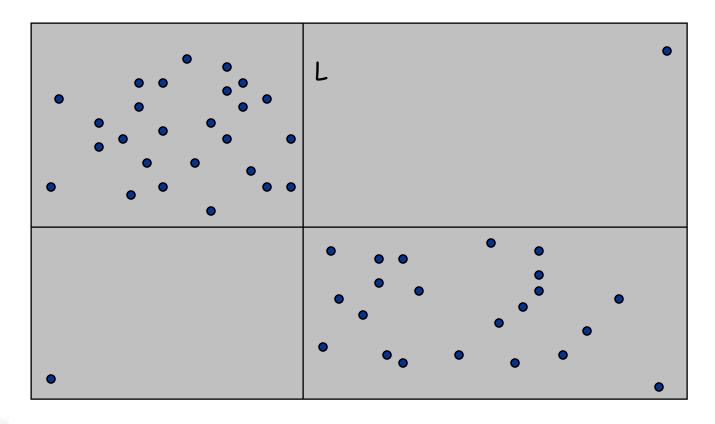


Closest Pair of Points: First Attempt

A STATE OF THE STA

Divide. Sub-divide region into 4 quadrants.

Obstacle. Impossible to ensure n/4 points in each piece.

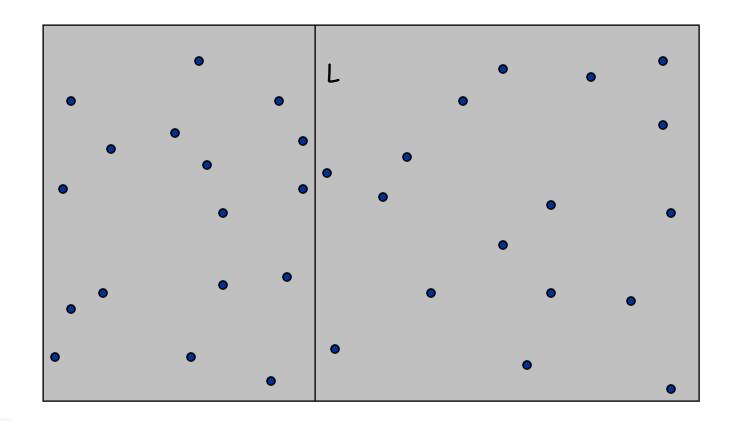






Algorithm.

Divide: draw vertical line L so that roughly ½n points on each side.

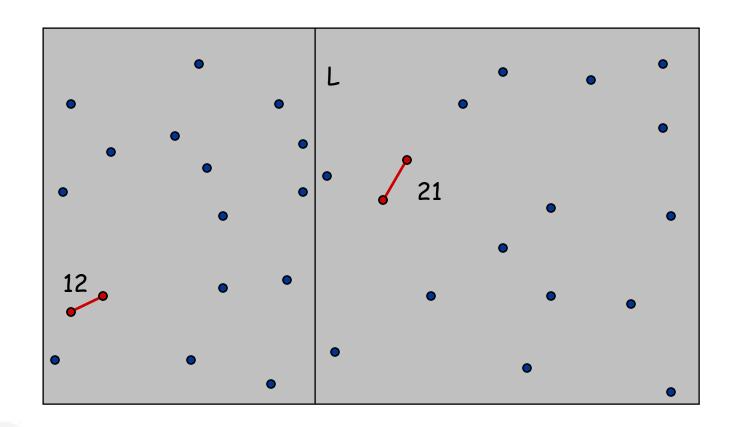






Algorithm.

- Divide: draw vertical line L so that roughly ½n points on each side.
- Conquer: find closest pair in each side recursively.

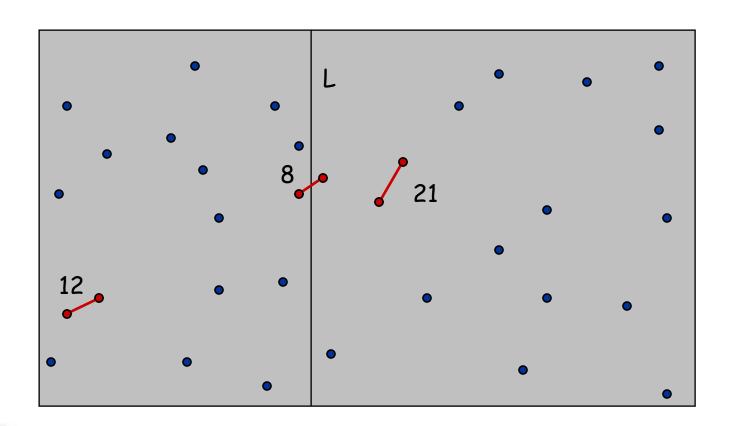






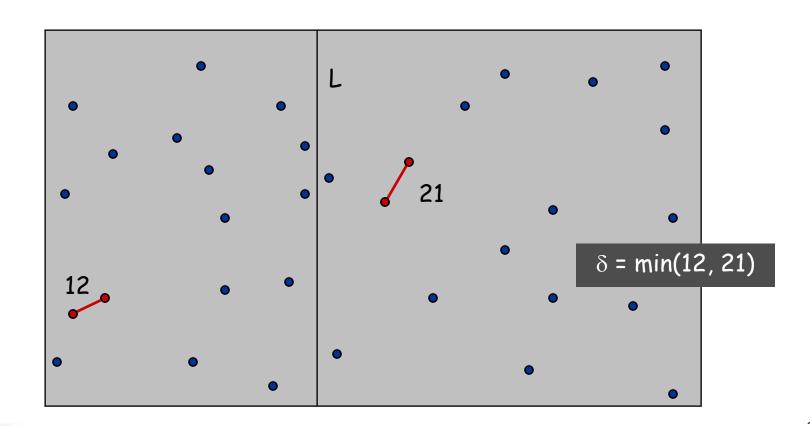
Algorithm.

- Divide: draw vertical line L so that roughly ½n points on each side.
- Conquer: find closest pair in each side recursively.
- _n Combine: find closest pair with one point in each side. ← seems like Θ(n²)
- Return best of 3 solutions.



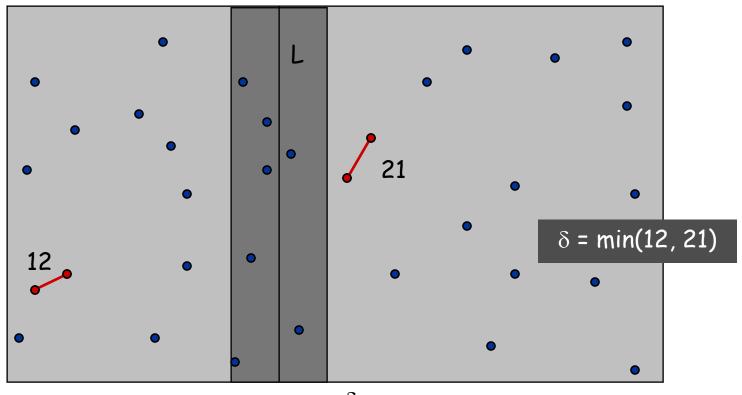
TO THE REPORT OF THE PARTY OF T

Find closest pair with one point in each side, assuming that distance $< \delta$.



Find closest pair with one point in each side, assuming that distance $< \delta$.

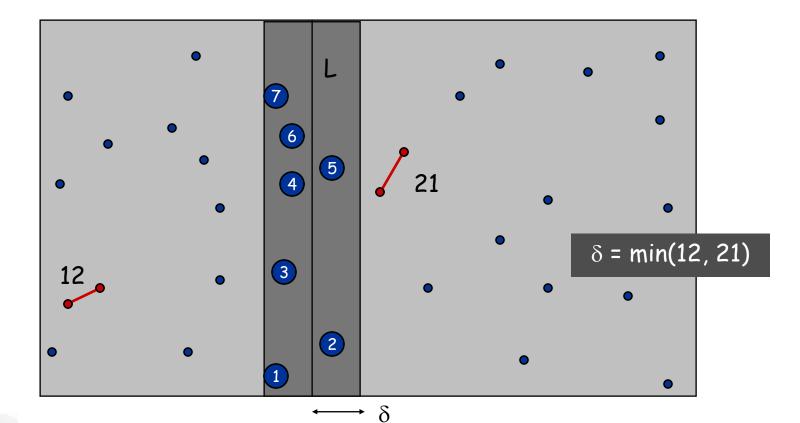
Observation: only need to consider points within δ of line L.





Find closest pair with one point in each side, assuming that distance $< \delta$.

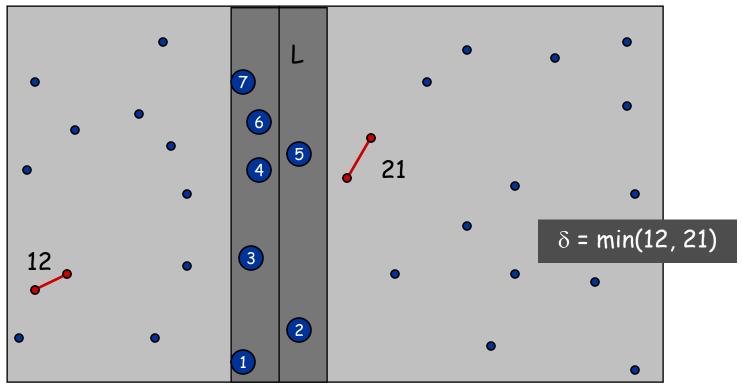
- . Observation: only need to consider points within δ of line L.
- Sort points in 2δ -strip by their y coordinate.





Find closest pair with one point in each side, assuming that distance $< \delta$.

- $_{\text{\tiny h}}$ Observation: only need to consider points within δ of line L.
- Sort points in 2δ -strip by their y coordinate.
- Only check distances of those within 11 positions in sorted list!





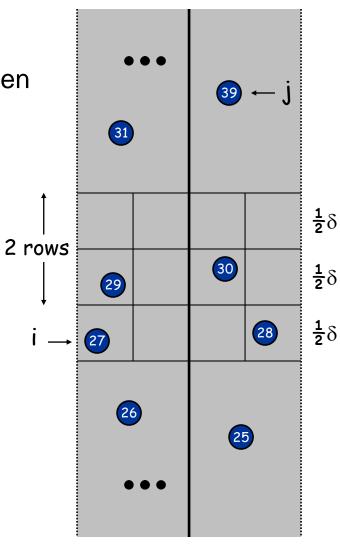
TO THE REAL PROPERTY OF THE PR

Def. Let s_i be the point in the 2δ -strip, with the i^{th} smallest y-coordinate.

Claim. If $|i - j| \ge 12$, then the distance between s_i and s_j is at least δ .

- No two points lie in same $\frac{1}{2}\delta$ -by- $\frac{1}{2}\delta$ box.
- Two points at least 2 rows apart have distance $\geq 2(\frac{1}{2}\delta)$.

Fact. Still true if we replace 12 with 7.



Closest Pair Algorithm



```
Closest-Pair (p_1, ..., p_n) {
   Compute separation line L such that half the points
                                                                       O(n \log n)
   are on one side and half on the other side.
   \delta_1 = Closest-Pair(left half)
                                                                       2T(n / 2)
   \delta_2 = Closest-Pair(right half)
   \delta = \min(\delta_1, \delta_2)
   Delete all points further than \delta from separation line L
                                                                       O(n)
                                                                        O(n \log n)
   Sort remaining points by y-coordinate.
   Scan points in y-order and compare distance between
                                                                        O(n)
   each point and next 11 neighbors. If any of these
   distances is less than \delta, update \delta.
   return \delta.
```

Closest Pair of Points: Analysis



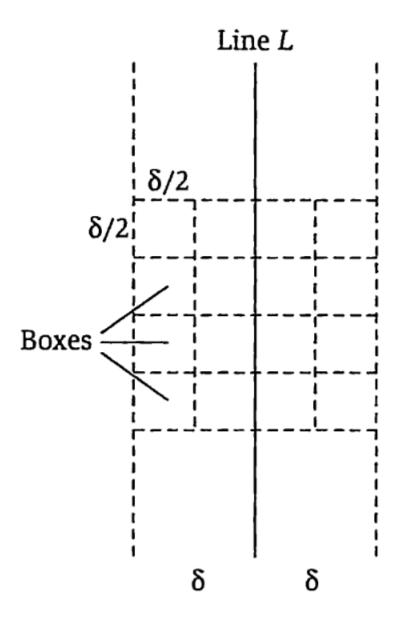
Running time.

$$T(n) \le 2T(n/2) + O(n \log n) \Rightarrow T(n) = O(n \log^2 n)$$

- Q. Can we achieve O(n log n)?
- A. Yes. Don't sort points in strip from scratch each time.
 - Each recursive returns two lists: all points sorted by y coordinate, and all points sorted by x coordinate.
 - Sort by merging two pre-sorted lists.

$$T(n) \le 2T(n/2) + O(n) \implies T(n) = O(n \log n)$$

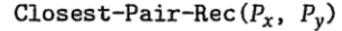




Closest-Pair(P)

Construct
$$P_x$$
 and P_y ($O(n \log n)$ time)

$$(p_0^*, p_1^*) = \text{Closest-Pair-Rec}(P_x, P_y)$$



If $|P| \leq 3$ then

find closest pair by measuring all pairwise distances Endif

Construct Q_x , Q_y , R_x , R_y (O(n) time)

$$(q_0^*, q_1^*) = \text{Closest-Pair-Rec}(Q_x, Q_y)$$

$$(r_0^*, r_1^*) = \text{Closest-Pair-Rec}(R_x, R_y)$$

$$\delta = \min(d(q_0^*, q_1^*), d(r_0^*, r_1^*))$$

 x^* = maximum x-coordinate of a point in set Q

$$L = \{(x,y) : x = x^*\}$$

 $S = \text{points in } P \text{ within distance } \delta \text{ of } L.$



Construct S_y (O(n) time)

For each point $s \in S_y$, compute distance from s to each of next 15 points in S_y Let s, s' be pair achieving minimum of these distances (O(n) time)

If $d(s,s') < \delta$ then Return (s,s')Else if $d(q_0^*,q_1^*) < d(r_0^*,r_1^*)$ then Return (q_0^*,q_1^*)

Else
Return (r_0^*, r_1^*) Endif





5.5 Integer Multiplication



Integer Arithmetic

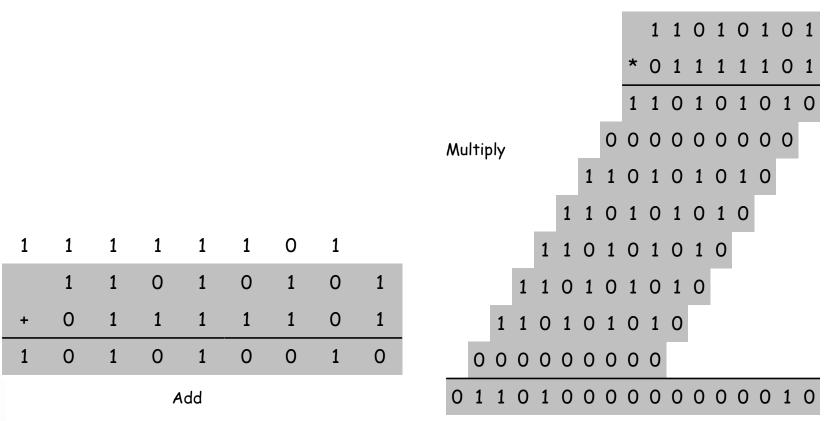


Add. Given two n-digit integers a and b, compute a + b.

O(n) bit operations.

Multiply. Given two n-digit integers a and b, compute a \times b.

Brute force solution: $\Theta(n^2)$ bit operations.



Divide-and-Conquer Multiplication: Warmup

To multiply two n-digit integers:

- Multiply four ½n-digit integers.
- Add two ½n-digit integers, and shift to obtain result.

$$x = 2^{n/2} \cdot x_1 + x_0$$

$$y = 2^{n/2} \cdot y_1 + y_0$$

$$xy = \left(2^{n/2} \cdot x_1 + x_0\right) \left(2^{n/2} \cdot y_1 + y_0\right) = 2^n \cdot x_1 y_1 + 2^{n/2} \cdot \left(x_1 y_0 + x_0 y_1\right) + x_0 y_0$$

$$T(n) = \underbrace{4T(n/2)}_{\text{recursive calls}} + \underbrace{\Theta(n)}_{\text{add, shift}} \Rightarrow T(n) = \Theta(n^2)$$

Î

assumes n is a power of 2

Karatsuba Multiplication



To multiply two n-digit integers:

- Add two ½n digit integers.
- Multiply three ½n-digit integers.
- Add, subtract, and shift ½n-digit integers to obtain result.

$$x = 2^{n/2} \cdot x_1 + x_0$$

$$y = 2^{n/2} \cdot y_1 + y_0$$

$$xy = 2^n \cdot x_1 y_1 + 2^{n/2} \cdot (x_1 y_0 + x_0 y_1) + x_0 y_0$$

$$= 2^n \cdot x_1 y_1 + 2^{n/2} \cdot ((x_1 + x_0)(y_1 + y_0) - x_1 y_1 - x_0 y_0) + x_0 y_0$$

$$A \qquad B \qquad A \qquad C \qquad C$$

Theorem. [Karatsuba-Ofman, 1962] Can multiply two n-digit integers in O(n^{1.585}) bit operations.

$$T(n) \leq \underbrace{T(\lfloor n/2 \rfloor) + T(\lfloor n/2 \rfloor) + T(1+ \lfloor n/2 \rfloor)}_{\text{recursive calls}} + \underbrace{\Theta(n)}_{\text{add, subtract, shift}}$$

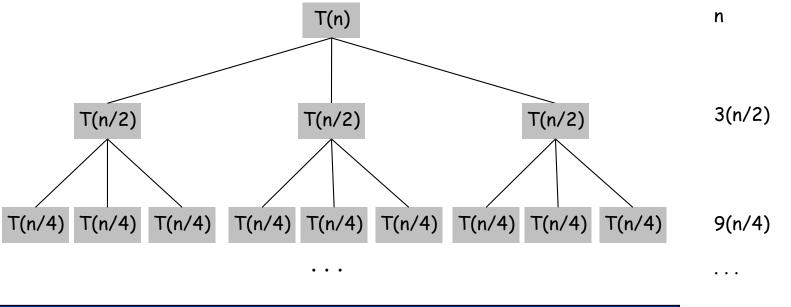
$$\Rightarrow T(n) = O(n^{\log_2 3}) = O(n^{1.585})$$

Karatsuba: Recursion Tree



$$T(n) = \begin{cases} 0 & \text{if } n = 1\\ 3T(n/2) + n & \text{otherwise} \end{cases}$$

$$T(n) = \sum_{k=0}^{\log_2 n} n \left(\frac{3}{2}\right)^k = \frac{\left(\frac{3}{2}\right)^{1 + \log_2 n} - 1}{\frac{3}{2} - 1} = 3n^{\log_2 3} - 2$$



••

T(2) T(2) T(2) T(2) T(2) T(2) T(2) T(2) T(2) T(2)



Recursive-Multiply(x,y):

Write
$$x = x_1 \cdot 2^{n/2} + x_0$$

 $y = y_1 \cdot 2^{n/2} + y_0$

Compute $x_1 + x_0$ and $y_1 + y_0$

 $p = \text{Recursive-Multiply}(x_1 + x_0, y_1 + y_0)$

 $x_1y_1 = \text{Recursive-Multiply}(x_1, y_1)$

 $x_0y_0 = \text{Recursive-Multiply}(x_0, y_0)$

Return $x_1y_1 \cdot 2^n + (p - x_1y_1 - x_0y_0) \cdot 2^{n/2} + x_0y_0$





Matrix Multiplication



Matrix Multiplication



Matrix multiplication. Given two n-by-n matrices A and B, compute C = AB.

$$c_{ij} = \sum_{k=1}^n a_{ik} b_{kj}$$

$$C_{ij} = \sum_{k=1}^{n} a_{ik} b_{kj}$$

$$\begin{bmatrix} c_{11} & c_{12} & \cdots & c_{1n} \\ c_{21} & c_{22} & \cdots & c_{2n} \\ \vdots & \vdots & \ddots & \vdots \\ c_{n1} & c_{n2} & \cdots & c_{nn} \end{bmatrix} = \begin{bmatrix} a_{11} & a_{12} & \cdots & a_{1n} \\ a_{21} & a_{22} & \cdots & a_{2n} \\ \vdots & \vdots & \ddots & \vdots \\ a_{n1} & a_{n2} & \cdots & a_{nn} \end{bmatrix} \times \begin{bmatrix} b_{11} & b_{12} & \cdots & b_{1n} \\ b_{21} & b_{22} & \cdots & b_{2n} \\ \vdots & \vdots & \ddots & \vdots \\ b_{n1} & b_{n2} & \cdots & b_{nn} \end{bmatrix}$$

Brute force. $\Theta(n^3)$ arithmetic operations.

Fundamental question. Can we improve upon brute force?



Matrix Multiplication: Warmup



Divide-and-conquer.

- Divide: partition A and B into ½n-by-½n blocks.
- Conquer: multiply 8 ½n-by-½n recursively.
- Combine: add appropriate products using 4 matrix additions.

$$\begin{bmatrix} C_{11} & C_{12} \\ C_{21} & C_{22} \end{bmatrix} = \begin{bmatrix} A_{11} & A_{12} \\ A_{21} & A_{22} \end{bmatrix} \times \begin{bmatrix} B_{11} & B_{12} \\ B_{21} & B_{22} \end{bmatrix}$$

$$C_{11} = (A_{11} \times B_{11}) + (A_{12} \times B_{21})$$

$$C_{12} = (A_{11} \times B_{12}) + (A_{12} \times B_{22})$$

$$C_{21} = (A_{21} \times B_{11}) + (A_{22} \times B_{21})$$

$$C_{22} = (A_{21} \times B_{12}) + (A_{22} \times B_{22})$$

$$T(n) = \underbrace{8T(n/2)}_{\text{recursive calls}} + \underbrace{\Theta(n^2)}_{\text{add, form submatrices}} \Rightarrow T(n) = \Theta(n^3)$$



Matrix Multiplication: Key Idea



Key idea. multiply 2-by-2 block matrices with only 7 multiplications.

$$\begin{bmatrix} C_{11} & C_{12} \\ C_{21} & C_{22} \end{bmatrix} = \begin{bmatrix} A_{11} & A_{12} \\ A_{21} & A_{22} \end{bmatrix} \times \begin{bmatrix} B_{11} & B_{12} \\ B_{21} & B_{22} \end{bmatrix} \qquad P_1 = A_{11} \times (B_{12} - B_{22})$$

$$C_{11} = P_5 + P_4 - P_2 + P_6$$

$$C_{12} = P_1 + P_2$$

$$C_{21} = P_3 + P_4$$

$$C_{22} = P_5 + P_1 - P_3 - P_7$$

$$P_{1} = A_{11} \times (B_{12} - B_{22})$$

$$P_{2} = (A_{11} + A_{12}) \times B_{22}$$

$$P_{3} = (A_{21} + A_{22}) \times B_{11}$$

$$P_{4} = A_{22} \times (B_{21} - B_{11})$$

$$P_{5} = (A_{11} + A_{22}) \times (B_{11} + B_{22})$$

$$P_{6} = (A_{12} - A_{22}) \times (B_{21} + B_{22})$$

$$P_{7} = (A_{11} - A_{21}) \times (B_{11} + B_{12})$$

- ⁿ 7 multiplications.
- $_{n}$ 18 = 10 + 8 additions (or subtractions).

Fast Matrix Multiplication



Fast matrix multiplication. (Strassen, 1969)

- Divide: partition A and B into ½n-by-½n blocks.
- Compute: 14 ½n-by-½n matrices via 10 matrix additions.
- Conquer: multiply 7 ½n-by-½n matrices recursively.
- Combine: 7 products into 4 terms using 8 matrix additions.

Analysis.

- Assume n is a power of 2.
- T(n) = # arithmetic operations.

$$T(n) = \underbrace{7T(n/2)}_{\text{recursive calls}} + \underbrace{\Theta(n^2)}_{\text{add, subtract}} \implies T(n) = \Theta(n^{\log_2 7}) = O(n^{2.81})$$



Fast Matrix Multiplication in Practice



Implementation issues.

- Sparsity.
- Caching effects.
- Numerical stability.
- Odd matrix dimensions.
- $_{n}$ Crossover to classical algorithm around n = 128.

Common misperception: "Strassen is only a theoretical curiosity."

- Advanced Computation Group at Apple Computer reports 8x speedup on G4 Velocity Engine when n ~ 2,500.
- Range of instances where it's useful is a subject of controversy.

Remark. Can "Strassenize" Ax=b, determinant, eigenvalues, and other matrix ops.



Fast Matrix Multiplication in Theory



- Q. Multiply two 2-by-2 matrices with only 7 scalar multiplications?
- A. Yes! [Strassen, 1969]

$$\Theta(n^{\log_2 7}) = O(n^{2.81})$$

- Q. Multiply two 2-by-2 matrices with only 6 scalar multiplications?
- A. Impossible. [Hopcroft and Kerr, 1971]

$$\Theta(n^{\log_2 6}) = O(n^{2.59})$$

- Q. Two 3-by-3 matrices with only 21 scalar multiplications?
- A. Also impossible.

$$\Theta(n^{\log_3 21}) = O(n^{2.77})$$

- Q. Two 70-by-70 matrices with only 143,640 scalar multiplications?
- A. Yes! [Pan, 1980]

$$\Theta(n^{\log_{70} 143640}) = O(n^{2.80})$$

Decimal wars.

- December, 1979: O(n^{2.521813}).
- January, 1980: O(n^{2.521801}).

Fast Matrix Multiplication in Theory



Best known. O(n^{2.376}) [Coppersmith-Winograd, 1987.]

Conjecture. $O(n^{2+\epsilon})$ for any $\epsilon > 0$.

Caveat. Theoretical improvements to Strassen are progressively less practical.

